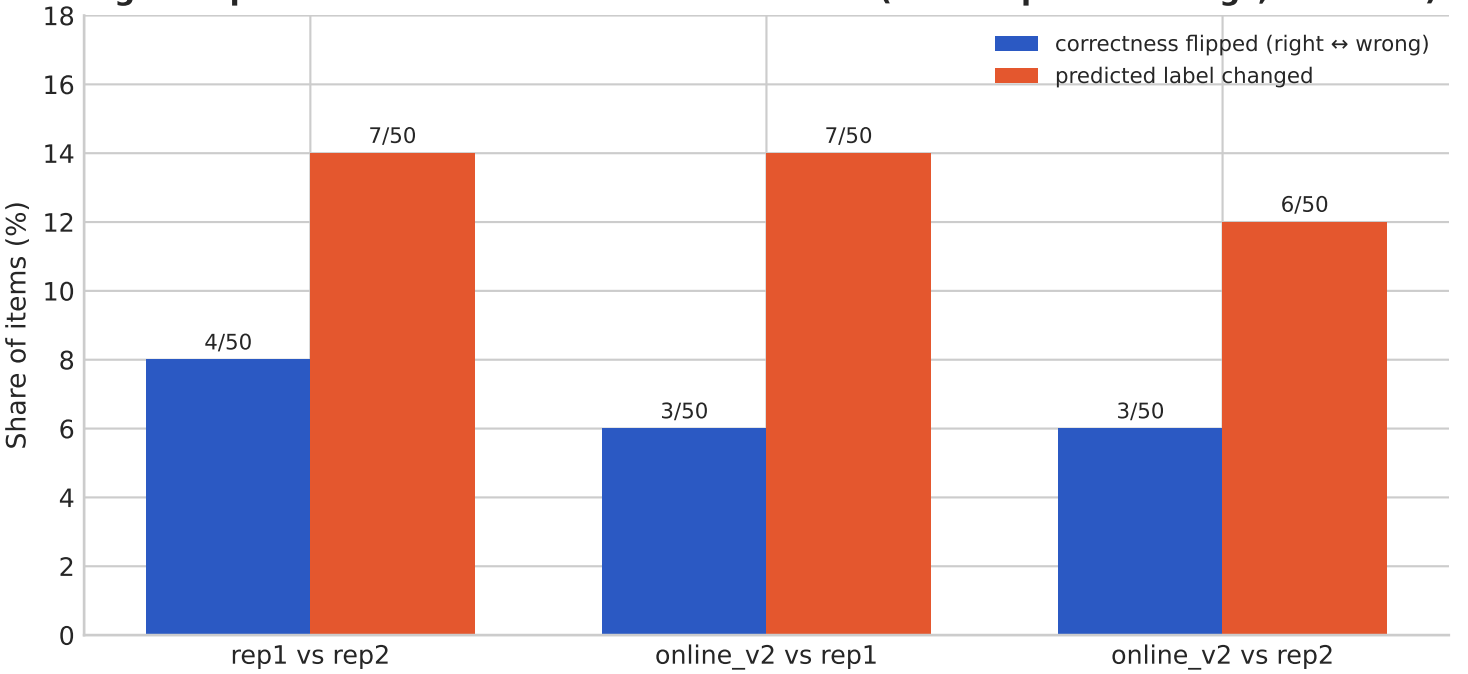
Fig 8. Repeat-call noise on the local LLM server (same input & settings, val n=50)
correctness flipped (right ↔ wrong)
predicted label changed
Share of items (%)
0
2
4
6
8
10
12
14
16
18
4/50
7/50
3/50
7/50
3/50
6/50
rep1 vs rep2
online_v2 vs rep1
online_v2 vs rep2